

Introduction to  
**High Performance Computing (HPC)**  
on Lewis and Clark clusters

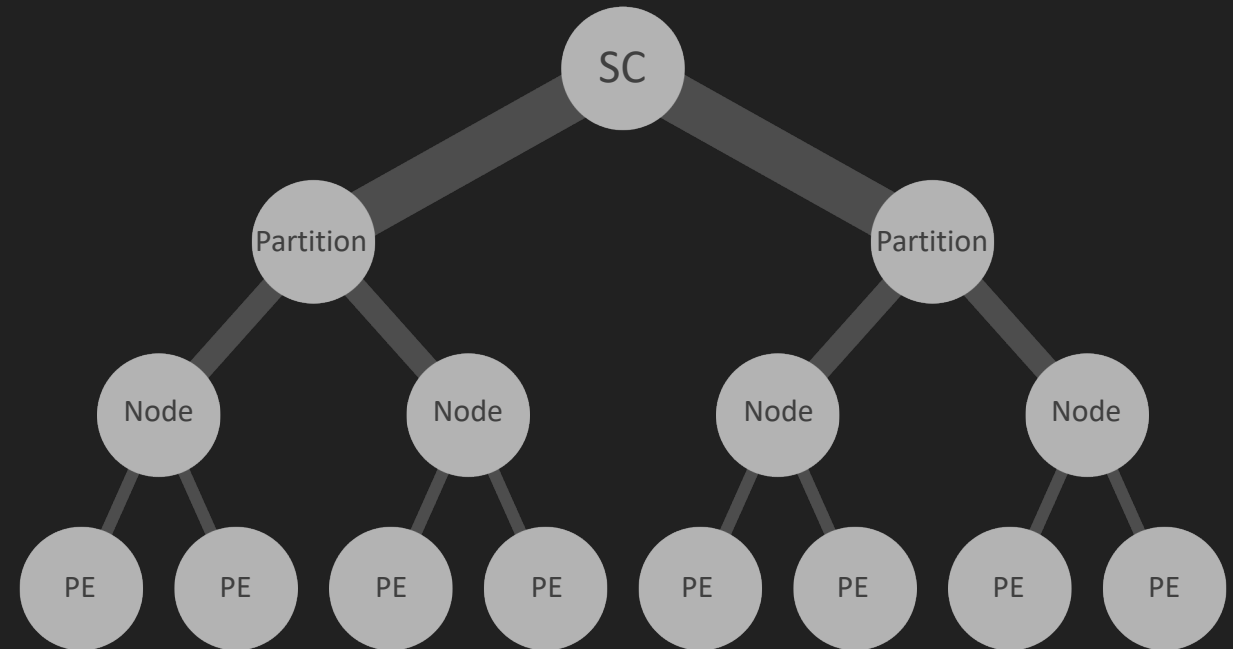
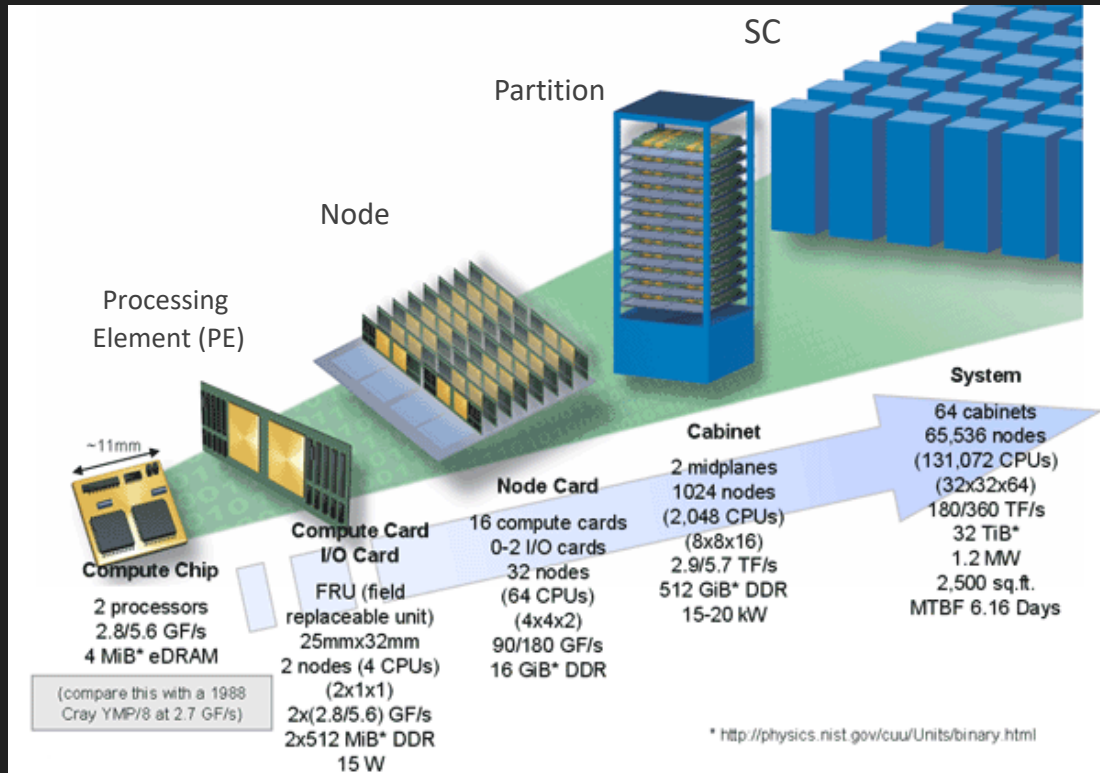
Research Computing Support Services (RCSS)  
<http://docs.rnet.missouri.edu>

RCSS CIE team

Ashkan Mirzaee, Asif Magdoom  
Brian Marxkors, Christina Roberts  
Predrag Lazic



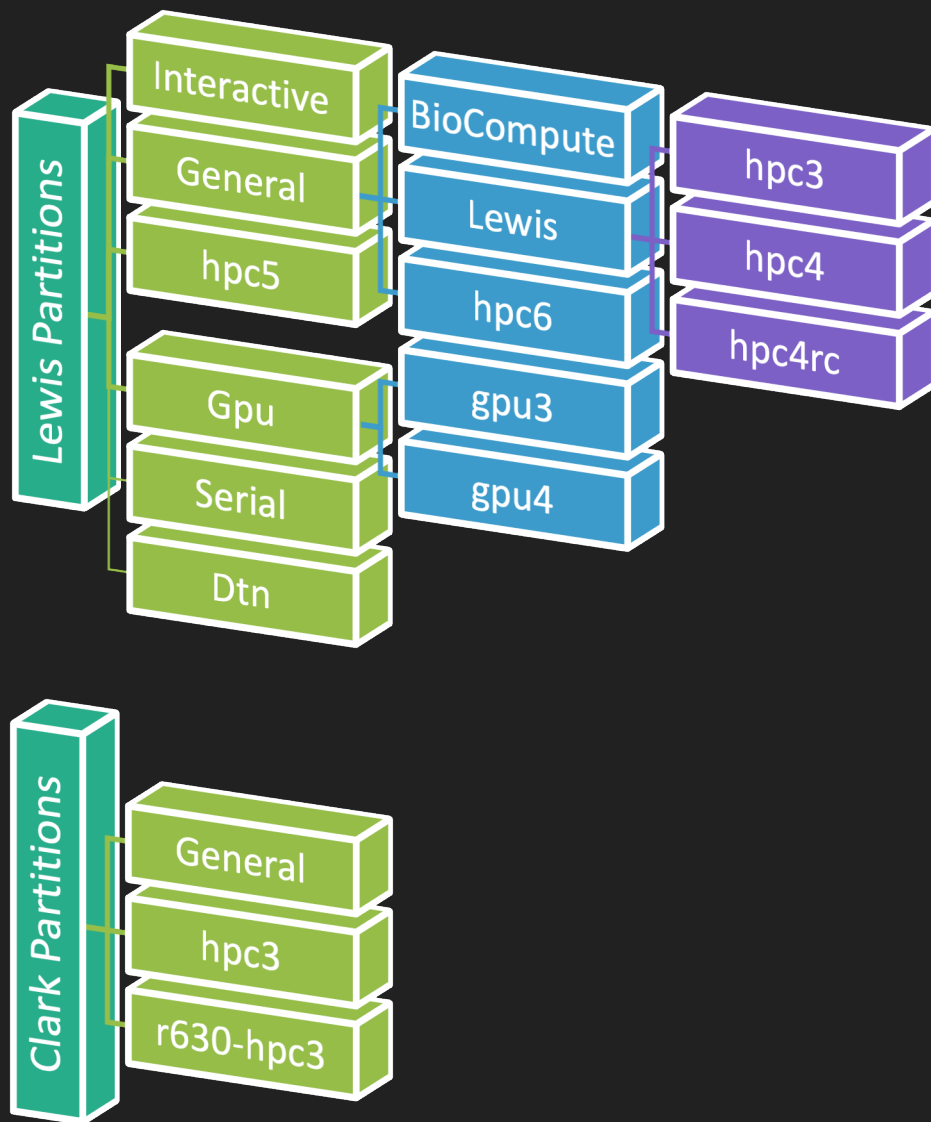
# What is a cluster (supercomputer)?



[wikimedia.org](https://www.wikipedia.org)

- A supercomputer is a computer with a **high level of performance** as compared to a general-purpose computer
- Purpose: **massive parallelization** because life is too short!
- The world's fastest 500 supercomputers run **Linux-based** operating systems

# Lewis and Clark



## Lewis

- A large-scale cluster for requesting high amount of resources
- Great for parallel programming
- GPU resources
- No cost for MU members for general usage
- Investment option is available to receive more resource (more fairshare)

## Clark

- Great for learning and teaching
- No need for registration and it is available to all MU members by MU username and PawPrint
- Usually less crowded - receive resources very fast
- No cost for MU members

# Lewis and Clark Partitions

Lewis

Partition Name	Time Limit	Nodes	Cores (per node*)	Cores (total)	Memory in GB (per nodes*)	Processors
Interactive	4:00:00	4	24+	144	251+	Intel(R) Xeon(R) CPU E5-2695 v2 @ 2.40GHz+
General	4:00:00	187	24+	5636	122+	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz+
BioCompute	2-00:00:00	37	56	2072	509	Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz
Lewis	2-00:00:00	90	24+	3564	122+	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz+
hpc3	2-00:00:00	22	24	1296	122+	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz
hpc4	2-00:00:00	37	28	1260	251+	Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz
hpc4rc	2-00:00:00	34	28	1008	251	Intel(R) Xeon(R) CPU E5-2680 v4 @ 2.40GHz
hpc6	2-00:00:00	61	48	2976	379+	Intel(R) Xeon(R) Gold 6252 CPU @ 2.10GHz
hpc5	2-00:00:00	35	40	1320	379	Intel(R) Xeon(R) Gold 6138 CPU @ 2.00GHz
Gpu	2:00:00	16	16+	372	122+	Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz+
gpu3	2-00:00:00	13	16+	284	122+	Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz
gpu4	2-00:00:00	3	40+	124	379+	Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz
Serial	2-00:00:00	1	64	64	1,025	AMD EPYC 7601 32-Core Processor
Dtn	2-00:00:00	2	16+	36	66+	Intel(R) Xeon(R) CPU X5550 @ 2.67GHz+

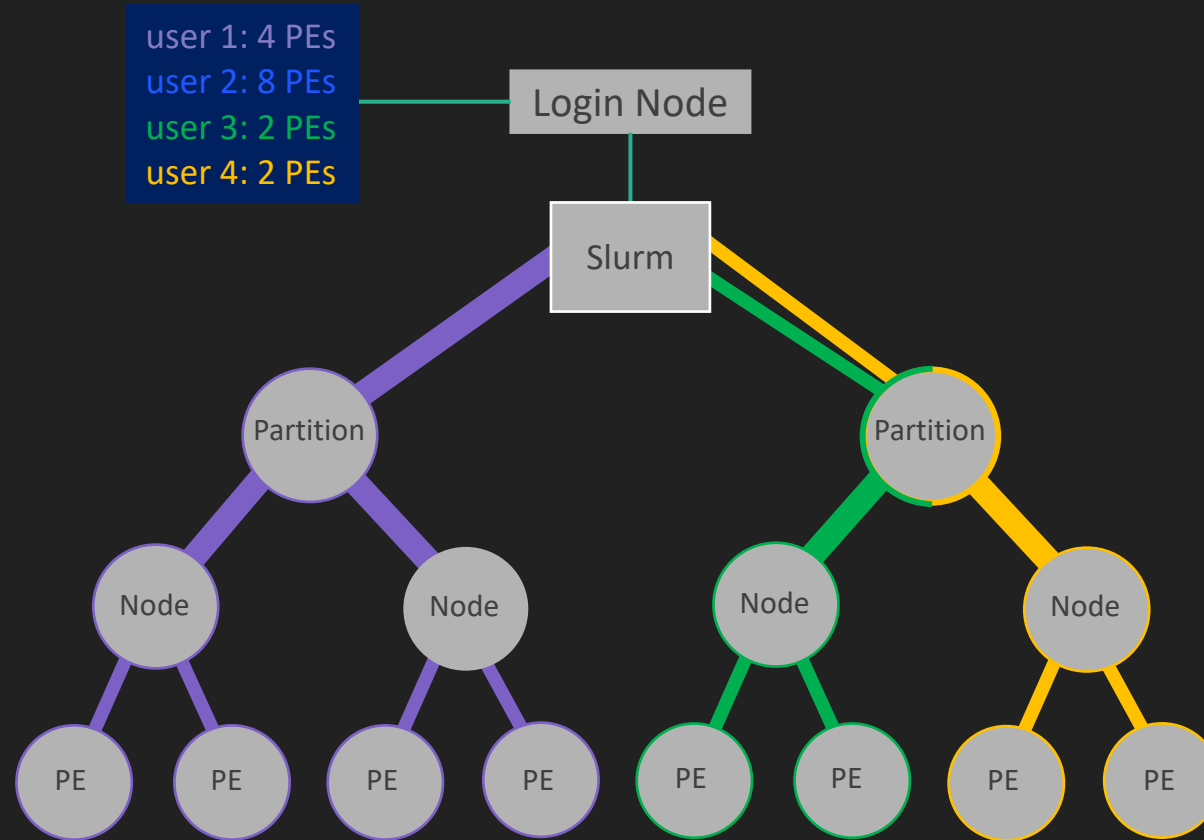
Clark

Partition Name	Time Limit	Nodes	Cores (per node*)	Cores (total)	Memory in MB (per nodes*)	Processors
General	2:00:00	4	24	96	122	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz
hpc3	2-00:00:00	4	24	96	122	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz
r630-hpc3	2-00:00:00	4	24	96	122	Intel(R) Xeon(R) CPU E5-2670 v3 @ 2.30GHz

\* plus sign (+) indicates a mixed environment. The number before the plus represents the minimum  
<http://docs.rnet.missouri.edu/policy/partition-policy/>

# SLURM

Slurm is a system for cluster management and job scheduling. All RCSS clusters use Slurm (<https://slurm.schedmd.com>).

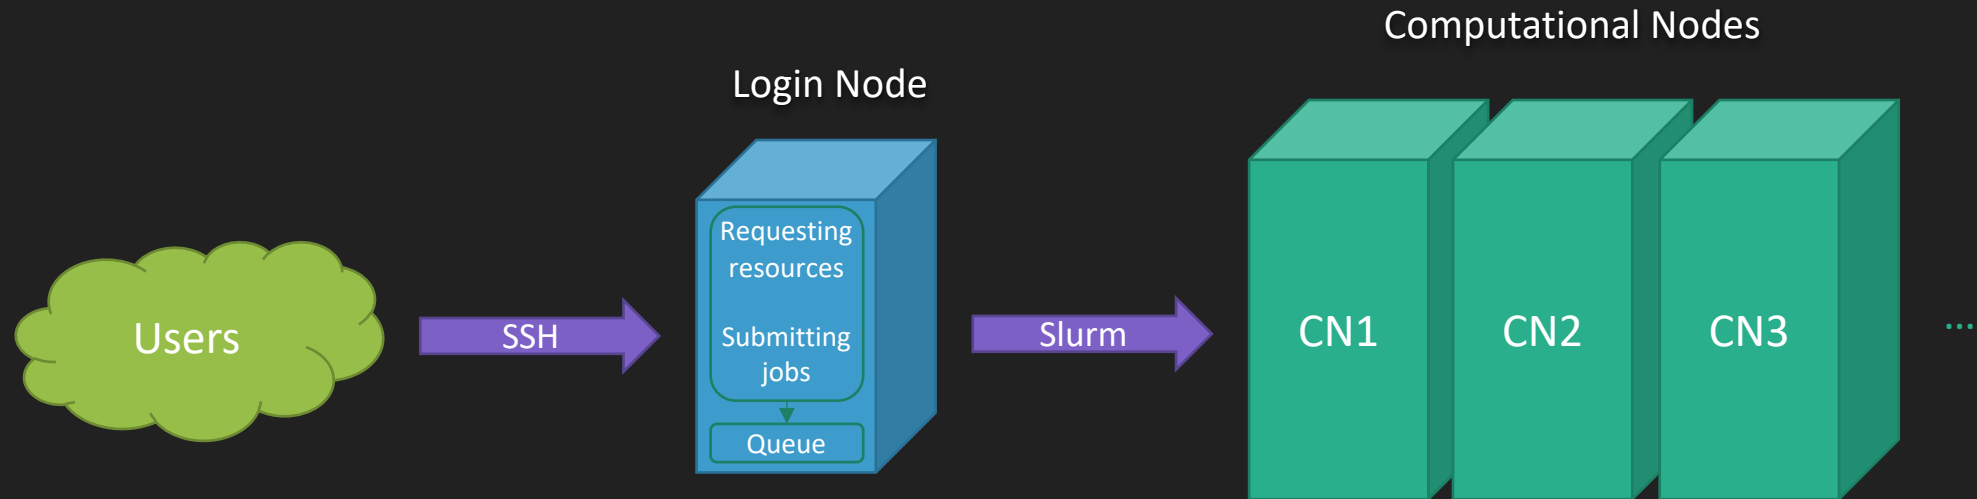


- Slurm is a **workload scheduler** and has set of tools for **submitting** and **monitoring** jobs
- Slurm is a **resource management system** and has many tools to find **available resources** in the cluster
- All Slurm commands start with **letter "s"**. In this course we will learn many of them
- Resource allocation depends on your fairshare i.e. priority in the queue

# Login node

All users connect to Clark and Lewis clusters through the login nodes.

```
[user@lewis4-r630-login-node675 ~]$  
[user@clark-r630-login-node907 ~]$
```



All jobs must be run using Slurm submitting tools to prevent running on the Lewis login node. Jobs that are found running on the login node will be immediately **terminated** followed up with a notification email to the user.

# Review

## Unix commands

man <code>manual</code>	hostname <code>host name</code>	cat <code>concatenate files and print output</code>
cd <code>change directory</code>	nproc <code>number of processers</code>	top <code>information about processes</code>
pwd <code>print working directory</code>	lscpu <code>list CPU architecture</code>	Tab <code>tab completion</code>
ls -la <code>long list of all files</code>	free -h <code>free memories</code>	clear <code>or Ctrl + l clear</code>
mkdir <code>make directory</code>	df <code>disk free</code>	exit <code>or Ctrl + d exit</code>
cp <code>copy</code>	echo \$USER <code>echo user id</code>	history <code>history</code>

## Text editors

<code>emacs -nw</code>	<code>nano -z</code>	<code>vim</code>
Ctrl-k <code>kill (cut)</code>	Ctrl-k <code>cut</code>	i <code>insert</code>
Ctrl-y <code>yank (paste)</code>	Ctrl-u <code>uncut (paste)</code>	Esc dd <code>delete line</code>
Alt-w <code>copy</code>	Alt-6 <code>copy</code>	Esc yy <code>yank (paste)</code>
Shift-arrow keys <code>select</code>	Alt-m + alt-a <code>select</code>	Esc u <code>undo</code> Esc p <code>paste</code>
Ctrl-z <code>suspend fg return</code>	Ctrl-z <code>suspend fg return</code>	Esc Ctrl + z <code>suspend fg ret.</code>
Ctrl-x + Ctrl-s <code>save</code>	Ctrl-o <code>write out</code>	Esc + :w <code>write</code>
Ctrl-x + Ctrl-c <code>close</code>	Ctrl-x <code>exit</code>	Esc + :q <code>quit :q! no change</code>

## Files

```
ssh username@lewis.rnet.missouri.edu connect to lewis by ssh  
cp -r /storage/hpc/data/plmx7/EXES_FOR_USERS/training ~ copy training files to the home directory
```

# Cluster info

**sinfo - sjstat - scontrol - sacctmgr**

**sinfo -s** summary of cluster resources -s --summarize  
**sinfo -p** <partitoion-name> -o %n,%C,%m,%z compute info of nodes in a partition -o --format  
**sinfo -p** Gpu -o %n,%C,%m,%G GPUs information in Gpu partition -p --partition  
**sjstat -c** show computing resources per node  
**scontrol** show partition <partition-name> partition information  
**scontrol** show node <node-name> node information  
**sacctmgr** show qos format=name,maxwall,maxsubmit show quality of services  
**./ncpu.py** show number of available CPUs and GPUs per node

```
* *****  
[amtwc@clark-r630-login-node907 ~]$ sinfo -s  
PARTITION AVAIL  TIMELIMIT  NODES(A/I/O/T)  NODELIST  
r630-hpc3    up 2-00:00:00      0/4/0/4  clark-r630-hpc3-node[908-911]  
hpc3        up 2-00:00:00      0/4/0/4  clark-r630-hpc3-node[908-911]  
General*    up   2:00:00      0/4/0/4  clark-r630-hpc3-node[908-911]  
[amtwc@clark-r630-login-node907 ~]$ sjstat -c
```

Scheduling pool data:

Pool	Memory	Cpus	Total	Usable	Free	Other	Traits
r630-hpc3	122534Mb	24	4	4	4		
hpc3	122534Mb	24	4	4	4		
General*	122534Mb	24	4	4	4		

CPUS/NODES(A/I/O/T) count of CPUs/nodes in the form "available/idle/other/total"

S:C:T counts number of "sockets, cores, threads"

The Scheduling data contains information pertaining to the:

Pool	a set of nodes
Memory	the amount of memory on each node
Cpus	the number of cpus on each node
Total	the total number of nodes in the pool
Usable	total usaable nodes in the pool
Free	total nodes that are currently free



# Users info

`sshare -U sacctmgr - df - lfs quota`

`sshare -U` show your fairshare and accounts

`-U --Users`

`sacctmgr show assoc user=$USER format=acc,user,share,qos,maxj` your QOS

`groups` show your groups

`df -h /home/$USER` home storage quota

`-h --human-readable`

`lfs quota -hg $USER /storage/hpc data/scratch` storage quota

`-g user/group`

`lfs quota -hg <group-name> /storage/hpc data/scratch` storage quota

`./userq.py` show user's fairshare, accounts, groups and QOS

- Resource allocation depends on your fairshare. If your fairshare is 0.000000 you have used the cluster more than your fair share and will be de-prioritized by the queuing software
- Users have 5GB at their home directory `/home/$USER` and 100GB at `/storage/hpc/data/$USER`
- **Do not** use home directory for running jobs, storing data or virtual environments
- Clark users have 100G on their home storage. The above methods does not apply for Clark
- The RCSS team reserves the right to delete anything in `/scratch` and `/local/scratch` at any time for any reason
- There are **no backups** of any storage. The RCSS team is not responsible for data integrity and data loss. **You are responsible for your own data and data backup**
- Review our storage policy at <http://docs.rnet.missouri.edu/policy/storage-policy/>

**Note:** `$USER` is a default environmental variable that returns your user id, try `echo $USER`

# Job submission

All jobs must be run using `srun` or `sbatch` to prevent running on the Lewis login node.

## `srun`: requesting resources to run jobs

```
srun <slurm-options> <software-name/path>
```

```
srun --pty /bin/bash
```

 requesting a pseudo terminal of bash shell to run jobs interactively

```
srun -p Interactive --pty /bin/bash
```

 requesting a p.t. of bash shell in Interactive Node -p --partition

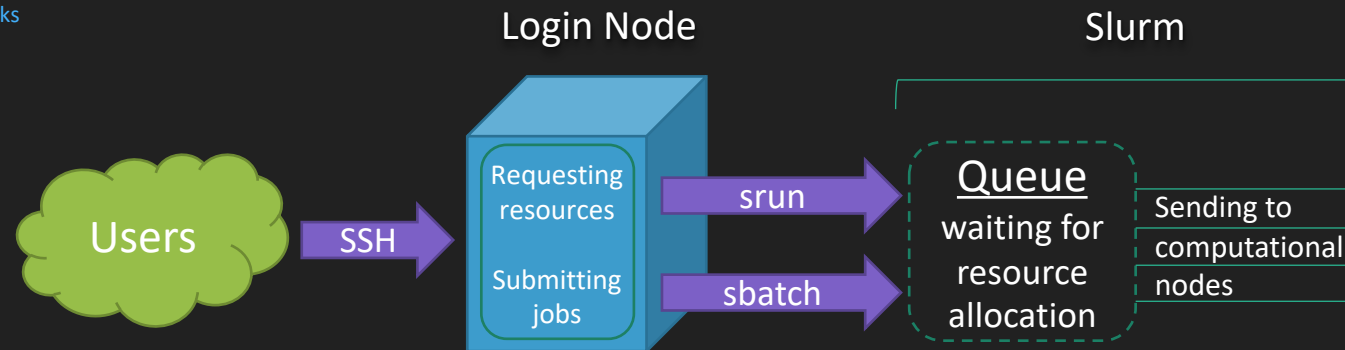
```
srun -p <partition-name> -n 4 --mem 16G --pty /bin/bash
```

 req. 4 tasks\* and 16G memory -n --ntasks

```
srun -p Gpu --gres gpu:1 -N 1 --ntasks-per-node 8 --pty /bin/bash
```

 req. 1 GPU and 1 node for running 8 tasks on Gpu partition -N --nodes

\*Slurm by default allocates 1 CPU per tasks



## `sbatch`: submitting jobs

Batch file is a shell script (`#!/bin/bash`) including Slurm options (`#SBATCH`) and computational tasks

```
sbatch <batch-file>
```

 submitting a batch file

After job completion we will receive outputs (`slurm-jobid.out`)

# Slurm options

## man srun - man sbatch

-p --partition <partition-name>	--pty <software-name/path>
--mem <memory>	--gres <general-resources>
-n --ntasks <number of tasks>	-t --time <days-hours:minutes>
-N --nodes <number-of-nodes>	-A --account <account>
-c --cpus-per-task <number-of-cpus>	-L --licenses <license>
-w --nodelist <list-of-node-names>	-J --job-name <jobname>

## Environmental Variables

\$SLURM\_JOB\_ID  
\$SLURM\_JOB\_NAME  
\$SLURM\_JOB\_NODELIST  
\$SLURM\_CPUS\_ON\_NODE  
\$SLURM\_SUBMIT\_HOST  
\$SLURM\_SUBMIT\_DIR

## Example

### test.py

```
#!/usr/bin/python3

import os

os.system("""
echo hostname: $(hostname)
echo number of processors: $(nproc)
echo data: $(date)
echo job id: $SLURM_JOB_ID
echo submit dir: $SLURM_SUBMIT_DIR
""")

print("Hello world")
```

### jobpy.sh

```
#!/bin/bash

#SBATCH -p Interactive
#SBATCH -n 4
#SBATCH --mem 8G

python3 test.py
```

### sbatch

```
sbatch jobpy.sh
```

### srun

```
srun -p Interactive -n 4 --mem 8G --pty bash

python3 test.py
```

### Output

```
hostname: lewis4-lenovo-hpc2-node282
number of processors: 4
data: Sun Jun 28 13:27:39 CDT 2020
job id: 21437062
submit dir: /home/user/training
Hello world
```

# Monitor jobs

**sacct -squeue -scancel**

sacct -X show your jobs in the last 24 hours -X --allocations  
sacct -X -S <yyyy-mm-dd> show your jobs since a date -S --starttime  
sacct -X -S <yyyy-mm-dd> -E <yyyy-mm-dd> -s <R/PD/F/CA/CG/CD> show running/pending/failed/cancelled/completing/completed jobs in a preiod of time -s --state  
sacct -j <jobid> show info about the jobid  
squeue -u <username> show a user jobs (R/PD/CD) in the queue -u --user  
squeue -u <username> --start show estimation time to start pending jobs  
scancel <jobid> cancel jobs  
./jobstat.py <day/week/month/year> show info about running, pending and completed jobs of a user within a time period (default is a week)

```
[amtwc@lewis4-r630-login-node675 alias]$ ./jobstat.py
```

Completed jobs for the last week:

JobID	User	Account	State	Partition	QOS	NCPUS	NNode	ReqMe	Submit	Reserved	Start	Elapsed	End	NodeList	JobName
21540719	amtwc	general	COMPLETED	Interact+	normal	1	1	8Gn	2020-07-07T12:28:40	00:00:00	2020-07-07T12:28:40	00:12:35	2020-07-07T12:41:15	lewis4-lenovo-hpc2-node282	bash
21541075	amtwc	general	COMPLETED	Interact+	normal	1	1	24Gn	2020-07-07T12:41:21	00:00:00	2020-07-07T12:41:21	00:16:13	2020-07-07T12:57:34	lewis4-lenovo-hpc2-node282	bash
21544202	amtwc	general	CANCELLED+	General	normal	1	1	1Gc	2020-07-08T10:09:53	00:00:06	2020-07-08T10:09:59	00:00:00	2020-07-08T10:09:59	None assigned	bash
21544203	amtwc	general	COMPLETED	Lewis	normal	1	1	1Gc	2020-07-08T10:10:10	00:00:00	2020-07-08T10:10:10	00:15:02	2020-07-08T10:25:12	lewis4-r630-hpc4-node674	bash
21544558	amtwc	general	COMPLETED	Interact+	normal	1	1	8Gn	2020-07-08T15:00:07	00:00:00	2020-07-08T15:00:07	00:26:03	2020-07-08T15:26:10	lewis4-lenovo-hpc2-node282	bash
21563592	amtwc	general	CANCELLED+	Lewis	normal	16	1	64Gn	2020-07-11T11:19:10	00:00:12	2020-07-11T11:19:22	00:00:00	2020-07-11T11:19:22	None assigned	bash

# Monitor CPU and Memory

## Completed jobs

### sacct - seff

```
sacct -j <jobid> -o User,Acc,AllocCPUS,Elaps,CPUTime,TotalCPU,AveDiskRead,AveDiskWrite,ReqMem,MaxRSS -j --jobs
info about CPU and virtual memory for completed jobs -o --format

seff <jobid> show jobs CPU and memory efficiency
```

```
[amtwc@lewis4-r630-login-node675 ~]$ sacct -j 20785018 -o User,Acc,AllocCPUS,Elaps,CPUTime,TotalCPU,AveDiskRead,AveDiskWrite,ReqMem,MaxRSS
User      Account AllocCPUS Elapsed CPUTime TotalCPU AveDiskRead AveDiskWrite ReqMem MaxRSS
-----
amtwc     general 16      00:48:39 12:58:24 01:49.774 66.58M      44.75M      64Gn    216K
```

```
[amtwc@lewis4-r630-login-node675 ~]$ seff 20785018
Job ID: 20785018
Cluster: lewis4
User/Group: amtwc/amtwc
State: COMPLETED (exit code 0)
Nodes: 1
```

Cores per node: 16

CPU Utilized: 00:01:50

CPU Efficiency: 0.24% of 12:58:24 core-walltime

Memory Utilized: 3.38 MB (estimated maximum)

Memory Efficiency: 0.01% of 64.00 GB (64.00 GB/node)

CPU Efficiency = CPU Utilization / Core-walltime  
Core-walltime = Core per node \* Elapsed time

## Running jobs:

### sstat - srun

```
sstat <jobid> -o AveCPU,AveDiskRead,AveDiskWrite,MaxRSS info about CPU and memory for running jobs (srun only)
```

```
srun --jobid <jobid> --pty /bin/bash attach to a srun/sbatch session and run `top` command to see information about processes
```

## Test

```
emacs -nw test.py
emacs -nw jobpy.sh
Ctrl+x+s to save and Ctrl+x+c to exit
sbatch jobpy.sh
sstat <jobid>
srun --jobid <jobid> --pty bash
top -u $USER press q to exit
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
13635	amtwc	20	0	160928	2916	1552	R	1.0	0.0	0:00.41	top
13435	amtwc	20	0	113284	1188	1020	S	0.0	0.0	0:00.00	slurm_script
13436	amtwc	20	0	233232	4752	1812	S	0.0	0.0	0:00.05	srun
13437	amtwc	20	0	28480	740	12	S	0.0	0.0	0:00.00	srun
13449	amtwc	20	0	124924	5612	2600	S	0.0	0.0	0:00.04	python3
13450	amtwc	20	0	124924	5612	2600	S	0.0	0.0	0:00.03	python3
13451	amtwc	20	0	124924	5612	2600	S	0.0	0.0	0:00.03	python3
13452	amtwc	20	0	124924	5612	2600	S	0.0	0.0	0:00.03	python3
13456	amtwc	20	0	108056	348	280	S	0.0	0.0	0:00.00	sleep
13457	amtwc	20	0	108056	348	280	S	0.0	0.0	0:00.00	sleep
13458	amtwc	20	0	108056	348	280	S	0.0	0.0	0:00.00	sleep
13464	amtwc	20	0	108056	348	280	S	0.0	0.0	0:00.00	sleep
13534	amtwc	20	0	115484	3868	1640	S	0.0	0.0	0:00.08	bash

resident set size (RES) = memory KB

#SBATCH -n 4  
#SBATCH --mem 8G

We are using 0% CPU and less than 6MB memory

# Modules

`module avail/load/unload/list/show/purge`

`module avail` available modules

`module show` show modules info

`module list` list loaded modules

`module load` loaded modules

`module unload` unload loaded modules

`module purge` unload all loaded modules

For example let's run R and MATLAB interactively:

## R

```
srun -p Interactive --mem 4G --pty /bin/bash
```

```
module load R
```

```
module list
```

```
1) R/R-3.3.3
```

```
R
```

## MATLAB

```
srun -p Interactive --mem 4G -L matlab --pty /bin/bash
```

```
module load matlab
```

```
module list
```

```
1) matlab/matlab-R2020a
```

```
matlab -nodisplay
```

**Never load modules in the login node.** It makes login node slow for all users. **Many modules don't work in the login node.** We can load modules in batch files, for example:

## test.R

```
#!/usr/bin/R
```

```
for (i in 1:3) {  
  cat("Hello world", i, "\n")  
}
```

## runr.sh

```
#!/bin/bash
```

```
#SBATCH -p Interactive  
#SBATCH --mem 4G
```

```
module load R  
Rscript test.R
```

## sbatch

```
sbatch runr.sh
```

## Output

```
Hello world 1  
Hello world 2  
Hello world 3
```

# What is next:

## Version control

- Git <https://git-scm.com/book/en/v2>

## Job dependencies

- Slurm dependency option (--dependency) <https://slurm.schedmd.com/sbatch.html>
- Snakemake <https://snakemake.readthedocs.io/>

## Virtual Environments

- Anaconda <https://conda.io/en/latest/>

## Software Installation

- Spack <https://spack.readthedocs.io/en/latest>

## Parallel programming

- Scientific languages C/Fortran
  - OpenMP
  - OpenACC for GPU parallelization
  - MPI for massive parallelization
- Python
  - NumPy
  - Numba
  - mpi4py

## RCSS Documentation

<http://docs.rnet.missouri.edu>

## XSEDE

<https://www.xsede.org/for-users/training>

## Software Carpentry

<https://software-carpentry.org/lessons/>

## HPC Carpentry

<https://hpc-carpentry.github.io>

## Data Carpentry

<https://datacarpentry.org/lessons/>

## Cornell Virtual Workshop

<https://cvw.cac.cornell.edu/topics>

## Pittsburgh Supercomputing Center (PSC)

<https://www.psc.edu/resources-for-users/training/>

## TACC Learning Portal

<https://learn.tacc.utexas.edu/course/>

## Feedback and Questions

[mudoitrcss@missouri.edu](mailto:mudoitrcss@missouri.edu)